Abdelbacet Mhamdi

Dr.-Ing. in Electrical Engineering

Senior Lecturer at ISET Bizerte

abdelbacet.mhamdi@bizerte.r-iset.tn

Machine Learning

LAB MANUAL

Institut Supérieur des Études Technologiques de Bizerte

Honor code

"During this course, you will be working with one or more partners with whom you may discuss any points concerning laboratory work. However, you must write your lab report, in your own words.

Lab reports that contain identical language are not acceptable, so do not copy your lab partner's writing.

If there is a problem with your data, include an explanation in your report. Recognition of a mistake and a well-reasoned explanation is more important than having high-quality data, and will be rewarded accordingly by your instructor. A lab report containing data that is inconsistent with the original data sheet will be considered a violation of the Honor Code.

Falsification of data or plagiarism of a report will result in prosecution of the offender(s) under the University Honor Code.

On your first lab report you must write out the entire honor pledge:

The work presented in this report is my own, and the data was obtained by my lab partner and me during the lab period.

On future reports, you may simply write *"Laboratory Honor Pledge"* and sign your name."

Contents

In order to activate the virtual environment and launch **Jupyter Lab**, you need to proceed as follow

① Press simultaneously the keys ⌨CTRL ⌨ALT and ⌨T on the keyboard[1];

② Type `mlpy` in the console prompt line;

> ▣ **Terminal**
>
> student@isetbz:~$ mlpy

③ Finally hit the ⏎ key.

KEEP THE SYSTEM CONSOLE OPEN.

▼ Remark 1

You should be able to utilize **Python** *from within the notebook through:*

Jupyter Lab   *at* *http://localhost:2468*

Marimo   *at* *http://localhost:1357*

🛑 Please use one of the provided templates when preparing your lab assessments:

LaTeX `https://www.overleaf.com/read/pwgpyvcxcvym#9e34eb`

Typst `https://typst.app/universe/package/ailab-isetbz`

---

[1]If you prefer using Windows, a similar environment has been setup for you by pressing ⊞ & ⌨R . This will open the dialog box `Run`. In the command line, type `cmd`, and then use the ⏎ key to confirm. Next, type `mlpy` and press ⏎ once more.

1 | *Python* **Onramp**

Student's name	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
Score /20	. . . .	. . . .	. . . .

Detailed Credits

Anticipation *(4 points)*	. . . .	. . . .	. . . .
Management *(2 points)*	. . . .	. . . .	. . . .
Testing *(7 points)*	. . . .	. . . .	. . . .
Data Logging *(3 points)*	. . . .	. . . .	. . . .
Interpretation *(4 points)*	. . . .	. . . .	. . . .

Motivations

★ *Python* is a popular programming language in the field of machine learning because it is relatively easy to learn and has a wide range of libraries and frameworks that support machine learning tasks.
★ *Python* has a large and active community of developers, which means that there are many resources available online, such as tutorials, documentation, and online forums, to help us learn and troubleshoot our code.
★ Many machine learning tools and frameworks, such as *TensorFlow* and *scikit-learn*, are written in *Python*, which makes it easy to integrate these tools into *Python* programs.
★ *Python* is a versatile language that can be used for a wide range of applications beyond machine learning, including web development, data analysis, and scientific computing.

The notebook is available at https://github.com/a-mhamdi/mlpy/ → *Codes* → *Python* → *Jupyter* → *py-onramp.ipynb*

Goals

1. Learn the basics of programming in *Python*;
2. Get familiar with *Jupyter Notebook*;
3. Use the modules of scientific computing.

1.1   Numerical variables & types

```
[1]: a = 1 # An integer
     print('The variable a = {} is of type {}'.format(a, type(a)))
```

```
The variable a = 1 is of type <class 'int'>
```

```
[2]: b = -1.25 # A floating number
     print('The variable b = {} is of type {}'.format(b, type(b)))
```

```
The variable b = -1.25 is of type <class 'float'>
```

```
[3]: c = 1+0.5j # A complex number
     print('The variable c = {} is of type {}'.format(c, type(c)))
```

```
The variable c = (1+0.5j) is of type <class 'complex'>
```

1.2   Strings

```
[4]: msg = "My 1st lab!"
     print(msg, type(msg), sep = '\n***\n') # \n: Carriage Return & Line Feed
     print(msg + 3* '\nPython is awesome')
```

```
My 1st lab!
***
<class 'str'>
My 1st lab!
Python is awesome
Python is awesome
Python is awesome
```

```
[5]: longMsg = """This is a long message,
     spanned over multiple lines"""
     print(longMsg)
```

```
This is a long message,
spanned over multiple lines
```

Indexing and slicing

```
[6]: # Positive indexing
     print(msg, msg[1:5], sep = ' -----> ')
     # Negative indexing
     print(msg, msg[-5:-1], sep = ' -----> ')
```

```
My 1st lab! -----> y 1s
My 1st lab! ----->  lab
```

String transformations

```
[7]: msg = 'A message'
     print(len(msg))
     print(msg.lower())
     print(msg.upper())
```

```
print(msg.split(' '))
print(msg.replace('mes', 'MES'))
print('a' in msg) #  Check if the variable `msg` contains the letter 'a'
```

```
9
a message
A MESSAGE
['A', 'message']
A MESsage
True
```

[8]:
```
price, number, perso = 300, 7, 'A customer'
print('{} asks for {} pieces. They cost {} TND!'.format(perso, number,␣
 ↪price))
print('{1} demande {2} pièces. They cost {0} TND!'.format(price, perso,␣
 ↪number))
```

```
A customer asks for 7 pieces. They cost 300 TND!
A customer demande 7 pièces. They cost 300 TND!
```

1.3   Binary, octal & hexadecimal

[9]:
```
x = 0b0101 # 0b : binary
print(x, type(x), sep = '\t----\t') # \t : tabular
y = 0xAF # 0x : hexadecimal
print(y, type(y), sep = '\t' + '---'*5 + '\t')
z = 0o010 # 0o : octal
print(z, type(z), sep = ', ')
```

```
5       ----    <class 'int'>
175     -------------- <class 'int'>
8, <class 'int'>
```

Boolean

[10]:
```
a = True
b = False
print(a)
print(b)
```

```
True
False
```

[11]:
```
print("50 > 20 ? : {} \n50 < 20 ? : {} \n50 = 20 ? : {}\n50 /= 20 ? : {}"
      .format(50 > 20, 50 < 20, 50 == 20, 50 != 20)
     )
```

```
50 > 20 ? : True
50 < 20 ? : False
50 = 20 ? : False
50 /= 20 ? : True
```

[12]:
```
print(bool(123), bool(0), bool('Lab'), bool())
```

True False True False

```
[13]: var1 = 100
      print(isinstance(var1, int))
      var2 = -100.35
      print(isinstance(var2, int))
      print(isinstance(var2, float))
```

```
True
False
True
```

1.4   Lists, tuples & dictionaries

In Python, a list is an ordered collection of items that can be of any data type (including other lists). Lists are defined using square brackets, with items separated by commas. For example:

```
[14]: shopping_list = ['milk', 'eggs', 'bread', 'apples']
```

A tuple is also an ordered collection of items, but it is immutable, meaning that the items it contains cannot be modified once the tuple is created. Tuples are defined using parentheses, with items separated by commas. For example:

```
[15]: point = (3, 5)
```

A dictionary is a collection of key-value pairs, where the keys are unique and used to look up the corresponding values. Dictionaries are defined using curly braces, with the key-value pairs separated by commas. The keys and values are separated by a colon. For example:

```
[16]: phonebook = {'Alice': '555-1234', 'Bob': '555-5678', 'Eve': '555-9101'}
```

You can access the items in a list or tuple using an index, and you can access the values in a dictionary using the corresponding keys. For example:

```
[17]: # Accessing the second item in a list
      print(shopping_list[1])  # prints 'eggs'

      # Accessing the first item in a tuple
      print(point[0])  # prints 3

      # Accessing the phone number for 'Bob' in the phonebook dictionary
      print(phonebook['Bob'])  # prints '555-5678'
```

```
eggs
3
555-5678
```

1.4.1   List

```
[18]: lst = ['a', 'b', 'c', 1, True] # An aggregate of various types
      print(lst)
```

```
['a', 'b', 'c', 1, True]
```

```
[19]: print(len(lst)) # Length of `lst` variable
      print(lst[1:3]) # Accessing elements of `lst`
      lst[0] = ['1', 0] # Combined list
      print(lst)
      print(lst[3:])
      print(lst[:3])
```

```
5
['b', 'c']
[['1', 0], 'b', 'c', 1, True]
[1, True]
[['1', 0], 'b', 'c']
```

```
[20]: lst.append('etc') # Insert 'etc' at the end
      print(lst)
```

```
[['1', 0], 'b', 'c', 1, True, 'etc']
```

```
[21]: lst.insert(1, 'xyz') # Inserting 'xyz'
      print(lst)
```

```
[['1', 0], 'xyz', 'b', 'c', 1, True, 'etc']
```

```
[22]: lst.pop(1)
      print(lst)
```

```
[['1', 0], 'b', 'c', 1, True, 'etc']
```

```
[23]: lst.pop()
      print(lst)
```

```
[['1', 0], 'b', 'c', 1, True]
```

```
[24]: del lst[0]
      print(lst)
```

```
['b', 'c', 1, True]
```

```
[25]: lst.append('b')
      print(lst)
      lst.remove('b')
      print(lst)
```

```
['b', 'c', 1, True, 'b']
['c', 1, True, 'b']
```

```
[26]: # Loop
      for k in lst:
          print(k)
```

```
c
1
True
b
```

```
[27]: lst.clear()
      print(lst)
```

```
[]
```

Method	Description
copy()	Returns a copy of the list
list()	Transforms into a list
extend ()	Extends a list by adding elements at its end
count()	Returns the occurrences of the specified value
index()	Returns the index of the first occurrence of a specified value
reverse()	Reverse a list
sort()	Sort a list

1.4.2 Tuples

```
[28]: tpl = (1, 2, 3)
      print(tpl)
```

```
(1, 2, 3)
```

```
[29]: tpl = (1, '1', 2, 'text')
      print(tpl)
```

```
(1, '1', 2, 'text')
```

```
[30]: print(len(tpl))
```

```
4
```

```
[31]: print(tpl[1:])
```

```
('1', 2, 'text')
```

```
[32]: try:
          tpl.append('xyz') # Throws an error
      except Exception as err:
          print(err)
```

```
'tuple' object has no attribute 'append'
```

```
[33]: try:
          tpl.insert(1, 'xyz') # Throws an error
      except Exception as err:
          print(err)
```

```
'tuple' object has no attribute 'insert'
```

```
[34]: my_lst = list(tpl)
      my_lst.append('xyz')
```

```python
print(my_lst, type(my_lst), sep = ', ')
```

```
[1, '1', 2, 'text', 'xyz'], <class 'list'>
```

```python
[35]: nv_tpl = tuple(my_lst) # Convert 'my_lst' into a tuple 'nv_tpl'
      print(nv_tpl, type(nv_tpl), sep = ', ')
```

```
(1, '1', 2, 'text', 'xyz'), <class 'tuple'>
```

```python
[36]: # Loop
      for k in nv_tpl:
          print(k)
```

```
1
1
2
text
xyz
```

```python
[37]: rs_tpl = tpl + nv_tpl
      print(rs_tpl)
```

```
(1, '1', 2, 'text', 1, '1', 2, 'text', 'xyz')
```

1.4.3 Dictionaries

```python
[38]: # dct = {"key": "value"}
      dct = {
          "Term" : "GM",
          "Speciality" : "ElnI",
          "Sem" : "4"
      }
      print(dct, type(dct), sep = ', ')
```

```
{'Term': 'GM', 'Speciality': 'ElnI', 'Sem': '4'}, <class 'dict'>
```

```python
[39]: print(dct["Sem"])
      sem = dct.get("Sem")
      print(sem)
```

```
4
4
```

```python
[40]: dct["Term"] = "GE"
      print(dct)
```

```
{'Term': 'GE', 'Speciality': 'ElnI', 'Sem': '4'}
```

```python
[41]: # Loop
      for el in dct:
          print(el, dct[el], sep = '\t|\t')
```

```
Term       |        GE
Speciality |        ElnI
Sem        |        4
```

```
[42]: for k in dct.keys():
          print(k)
```

```
Term
Speciality
Sem
```

```
[43]: for v in dct.values():
          print(v)
```

```
GE
ElnI
4
```

1.5   NumPy

NumPy is a *Python* library that is used for scientific computing and data analysis. It provides support for large, multi-dimensional arrays and matrices of numerical data, and a large library of mathematical functions to operate on these arrays.

One of the main features of *NumPy* is its *N*-dimensional array object, which is used to store and manipulate large arrays of homogeneous data (*i.e.*, data of the same type, such as integers or floating point values). The array object provides efficient operations for performing element-wise calculations, indexing, slicing, and reshaping.

NumPy also includes a number of functions for performing statistical and mathematical operations on arrays, such as mean, standard deviation, and dot product. It also includes functions for linear algebra, random number generation, and Fourier transforms.

Official documentation can be found at https://numpy.org/

```
[44]: import numpy as np
```

NumPy vs List

```
[45]: a_np = np.arange(6) # NumPy
      print("a_np = ", a_np)
      print(type(a_np))
      a_lst = list(range(0,6)) # List
      print("a_lst = ", a_lst)
      print(type(a_lst))
      # Comparison
      print("2 * a_np = ", a_np * 2)
      print("2 * a_lst = ", a_lst * 2)
```

```
a_np =  [0 1 2 3 4 5]
<class 'numpy.ndarray'>
a_lst =  [0, 1, 2, 3, 4, 5]
<class 'list'>
2 * a_np =  [ 0  2  4  6  8 10]
2 * a_lst =  [0, 1, 2, 3, 4, 5, 0, 1, 2, 3, 4, 5]
```

```
[46]: v_np = np.array([1, 2, 3, 4, 5, 6]) # NB : parentheses then brackets, i.e,␣
      ↪([])
```

```
print(v_np)
```

```
[1 2 3 4 5 6]
```

[47]:
```
v_np = np.array([[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12]])
print(v_np)
```

```
[[ 1  2  3  4]
 [ 5  6  7  8]
 [ 9 10 11 12]]
```

[48]:
```
print(type(v_np))
```

```
<class 'numpy.ndarray'>
```

[49]:
```
print(v_np[0])
```

```
[1 2 3 4]
```

[50]:
```
v_np.ndim # Dimensions of v_np
```

[50]: 2

[51]:
```
v_np.shape # Number of lignes and columns, may be more
```

[51]: (3, 4)

[52]:
```
v_np.size # How many elements are in `v_np`
```

[52]: 12

If we need to create a matrix $(3, 3)$, we can do as follows:

[53]:
```
u = np.arange(9).reshape(3,3)
print(u)
```

```
[[0 1 2]
 [3 4 5]
 [6 7 8]]
```

Let us see some known operations to do on matrices

[54]:
```
M = np.array([[1, 2], [1, 2]])
print(M)
```

```
[[1 2]
 [1 2]]
```

[55]:
```
N = np.array([[0, 3], [4, 5]])
print(N)
```

```
[[0 3]
 [4 5]]
```

Addition

```
[56]: print(M + N)
      print(np.add(M, N))
```

```
[[1 5]
 [5 7]]
[[1 5]
 [5 7]]
```

Subtraction

```
[57]: print(M-N)
      print(np.subtract(M, N))
```

```
[[ 1 -1]
 [-3 -3]]
[[ 1 -1]
 [-3 -3]]
```

Element-wise Division

$$\begin{bmatrix} 0 & 3 \\ 4 & 5 \end{bmatrix} ./ \begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} = \begin{bmatrix} 0:1 & 3:2 \\ 4:1 & 5:2 \end{bmatrix}$$

```
[58]: print(N / M)
      print(np.divide(N, M))
```

```
[[0.  1.5]
 [4.  2.5]]
[[0.  1.5]
 [4.  2.5]]
```

Element-wise Product

Element-wise multiplication, also known as **Hadamard product**, is an operation that multiplies each element of one matrix with the corresponding element of another matrix. It is denoted by the symbol $\odot$ or .* in some programming languages.

For example, consider the following matrices:

$$A = \begin{bmatrix} a_1, & a_2, & a_3 \end{bmatrix} \quad \text{and} \quad B = \begin{bmatrix} b_1, & b_2, & b_3 \end{bmatrix}$$

The element-wise product of these matrices is:

$$A \odot B = \begin{bmatrix} a_1 b_1, & a_2 b_2, & a_3 b_3 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} .\times \begin{bmatrix} 0 & 3 \\ 4 & 5 \end{bmatrix} = \begin{bmatrix} 0 & 6 \\ 4 & 10 \end{bmatrix}$$

We need element-wise multiplication in many applications. For example, in image processing, element-wise multiplication is used to modify the intensity values of an image by multiplying each pixel value with a scalar

value. In machine learning, element-wise multiplication is used in the implementation of various neural network layers, such as convolutional layers and fully connected layers. Element-wise multiplication is also used in many other mathematical and scientific applications.

```
[59]: print(M * N)
      print(np.multiply(M, N))
```

```
[[ 0  6]
 [ 4 10]]
[[ 0  6]
 [ 4 10]]
```

Dot Product

$$
\begin{bmatrix} 1 & 2 \\ 1 & 2 \end{bmatrix} \times \begin{bmatrix} 0 & 3 \\ 4 & 5 \end{bmatrix} = \begin{bmatrix} 8 & 13 \\ 8 & 13 \end{bmatrix}
$$

```
[60]: print(M.dot(N))
      print(np.dot(M, N))
```

```
[[ 8 13]
 [ 8 13]]
[[ 8 13]
 [ 8 13]]
```

Kronecker Product

$$
\begin{bmatrix} 1 & 2 & 3 & 4 \end{bmatrix} \bigotimes \begin{bmatrix} 1 & 2 \\ 3 & 4 \\ 5 & 6 \\ 7 & 8 \end{bmatrix} = \begin{bmatrix} 1 & 2 & 2 & 4 & 3 & 6 & 4 & 8 \\ 3 & 4 & 6 & 8 & 9 & 12 & 12 & 16 \\ 5 & 6 & 10 & 12 & 15 & 18 & 20 & 24 \\ 7 & 8 & 14 & 16 & 21 & 24 & 28 & 32 \end{bmatrix}
$$

```
[61]: u = np.arange(1, 5)
      v = np.arange(1, 9).reshape(4, 2)
      u, v
```

```
[61]: (array([1, 2, 3, 4]),
       array([[1, 2],
              [3, 4],
              [5, 6],
              [7, 8]]))
```

```
[62]: np.kron(u, v)
```

```
[62]: array([[ 1,  2,  2,  4,  3,  6,  4,  8],
             [ 3,  4,  6,  8,  9, 12, 12, 16],
             [ 5,  6, 10, 12, 15, 18, 20, 24],
             [ 7,  8, 14, 16, 21, 24, 28, 32]])
```

Determinant of a matrix

```
[63]: print("Determinant of M:")
      print(np.linalg.det(M))
      print("Determinant of N:")
      print(np.linalg.det(N))
```

```
Determinant of M:
0.0
Determinant of N:
-12.0
```

1.6 Matplotlib

Matplotlib is a 2D data visualization library in *Python* that allows users to create a wide range of static, animated, and interactive visualizations in *Python*. It is one of the most widely used data visualization libraries in the *Python* data science ecosystem and is particularly useful for creating line plots, scatter plots, bar plots, error bars, histograms, bar charts, pie charts, box plots, and many other types of visualizations.

Matplotlib is built on top of *NumPy* and is often used in conjunction with other libraries in the PyData ecosystem, such as *Pandas* and *Seaborn*, to create complex visualizations of data. It is also compatible with a number of different backends, such as the *Jupyter notebook*, *Qt*, and *Tkinter*, which allows it to be used in a wide range of environments and contexts.

The full documentation and an exhaustive list of samples can be found at https://matplotlib.org/

```
[64]: import numpy as np
      from matplotlib import pyplot as plt

      plt.style.use("ggplot")
      plt.rcParams['figure.figsize'] = [8, 4]
```

We begin by creating a sinusoidal waveform denoted by *x*, period is 1 sec. The offset is 1.

```
[65]: # Continuous function
      t = np.arange(0.0, 2.0, 0.01)
      x = 1 + np.sin(2 * np.pi * t) # Frequency = 1Hz
```

The set of instructions that allow to plot (x) are:

```
[66]: plt.plot(t, x)

      # Give the graph a title
      plt.title(r"$x(t) = 1+\sin\left(2\pi\frac{t}{1}\right)$")
      plt.xlabel("$t$ (sec)"); # Label the axis
```

$$x(t) = 1 + \sin\left(2\pi\frac{t}{1}\right)$$


```
[67]:  # Discret Function
       t = np.arange(0.0, 2.0, 0.1)
       y = np.sin(2*np.pi*t) # Same thing! Sinusoidal signal
```

```
[68]:  plt.stem(t, y)
       plt.xlabel("$t$ (sec)")
```

```
[68]:  Text(0.5, 0, '$t$ (sec)')
```



```
[69]:  x = np.logspace(-2, 3, 100)
       y = np.log10(x)
```

```
[70]:  np.log10?
```

```
[71]:  plt.plot(x, y)
```

[71]: [<matplotlib.lines.Line2D at 0x7f2199ca91e0>]

[72]:
```python
plt.semilogx(x, y)
```

[72]: [<matplotlib.lines.Line2D at 0x7f2157bb5a50>]

About distributions

[73]:
```python
a = np.random.randn(2**16) # Normal Distribution
b = np.random.rand(2**16)  # Uniform Distribution
```

[74]:
```python
_, ax = plt.subplots(1, 2)
ax[0].hist(a, bins=16)
ax[1].hist(b, bins=16);
```

```
[75]:  _, ax = plt.subplots(2, 2)
       ax[0][0].hist2d(a, a, bins=32)
       ax[0][1].hist2d(a, b, bins=32)
       ax[1][0].hist2d(b, a, bins=32)
       ax[1][1].hist2d(b, b, bins=32);
```

Task № 1:

In each instance, you are required to put in place a program that enables to

a) without using builtin functions

- take a string as input and returns the string with all vowels removed;
- determine whether a given year is a leap year or not;
- take a list of integers as input and returns the sum of all positive numbers in the list;
- take a string as input and checks if it's a palindrome. Return True if it is, False otherwise.

b) using *NumPy*

- generate random numbers and calculate the mean and standard deviation of these numbers;
- take a 1D *NumPy* array as input and returns the sum of squares of all elements;
- take a 1D *NumPy* array as input and returns the indices of all occurrences of the maximum value.

c) using *Matplotlib*

- create a simple plot of the sine function from 0 to 2π.
- visualize the distribution of a random dataset generated using *NumPy*;
- take a list of numbers as input and creates a histogram of these numbers;
- plot the first few terms of the Fibonacci sequence;
- take a string as input and creates a bar chart showing the frequency of each character in the string.

d) calculate electricity bills for regular low voltage customers in **Tunisia** according to **STEG**'s official tariff structure, using the official tariff information. The deliverable is a working console application that accurately calculates and displays electricity bills with detailed breakdowns for verification purposes.

2 | Linear Regression

Student's name	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
Score /20	. . . .	. . . .	. . . .

Detailed Credits

Anticipation **(4 points)**	. . . .	. . . .	. . . .
Management **(2 points)**	. . . .	. . . .	. . . .
Testing **(7 points)**	. . . .	. . . .	. . . .
Data Logging **(3 points)**	. . . .	. . . .	. . . .
Interpretation **(4 points)**	. . . .	. . . .	. . . .

> **Motivations**
>
> ★ Linear regression is a fundamental statistical technique that is widely used in many fields, including economics, finance, biology, and computer science. It is a simple and effective way to model the relationship between a dependent variable and one or more independent variables.
> ★ Linear regression is relatively easy to understand and implement, making it a good starting point for us who are new to statistical modeling. It is also a good foundation for learning more advanced statistical techniques, such as multivariate or logistic regression.
> ★ Linear regression can be an useful tool for making predictions and understanding the underlying trends in data. It can help us to better understand and analyze data, and to make informed decisions based on our findings.

The notebook is available at `https://github.com/a-mhamdi/mlpy/` → *Codes* → *Python* → *Jupyter* → *multiple-linear-regression.ipynb*

Multiple linear regression is a type of regression analysis in which there are multiple independent variables that have an effect on the dependent variable. In multiple linear regression, the goal is to find the linear equation that best explains the relationship between the outcome and the features in X.

The equation takes the form:

$$y = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_{m-1} x_{m-1}$$

where y is the dependent variable, $x_1, x_2, …, x_{m-1}$ are the independent variables, and $\theta_0, \theta_1, \theta_2, …, \theta_{m-1}$ are the coefficients that represent the influence of each variable on the output y. The coefficients are estimated

using the data, and the resulting equation can be used later to make predictions on new data.

Importing the libraries

```python
[1]: import numpy as np
     import pandas as pd
     from matplotlib import pyplot as plt
```

```python
[2]: np.set_printoptions(precision=3)
```

```python
[3]: # Show plots in an interactive format, e.g., zooming, saving, etc
     %matplotlib inline
```

```python
[4]: plt.style.use('ggplot')
```

Importing the dataset

```python
[5]: df = pd.read_csv('./Datasets/50_Startups.csv')
```

```python
[6]: df.head()
```

```
[6]:    R&D Spend  Administration  Marketing Spend       State     Profit
    0  165349.20       136897.80        471784.10    New York  192261.83
    1  162597.70       151377.59        443898.53  California  191792.06
    2  153441.51       101145.55        407934.54     Florida  191050.39
    3  144372.41       118671.85        383199.62    New York  182901.99
    4  142107.34        91391.77        366168.42     Florida  166187.94
```

```python
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 5 columns):
 #   Column           Non-Null Count  Dtype
---  ------           --------------  -----
 0   R&D Spend        50 non-null     float64
 1   Administration   50 non-null     float64
 2   Marketing Spend  50 non-null     float64
 3   State            50 non-null     object
 4   Profit           50 non-null     float64
dtypes: float64(4), object(1)
memory usage: 2.1+ KB
```

```python
[8]: df.describe()
```

```
[8]:            R&D Spend  Administration  Marketing Spend         Profit
    count      50.000000       50.000000        50.000000      50.000000
    mean    73721.615600   121344.639600    211025.097800  112012.639200
    std     45902.256482    28017.802755    122290.310726   40306.180338
    min         0.000000    51283.140000         0.000000   14681.400000
    25%     39936.370000   103730.875000    129300.132500   90138.902500
    50%     73051.080000   122699.795000    212716.240000  107978.190000
```

```
75%      101602.800000    144842.180000    299469.085000  139765.977500
max      165349.200000    182645.560000    471784.100000  192261.830000
```

Extract features *X* and target *y* from the dataset. **Profit** is the dependant variable.

```
[9]:  X = df.iloc[:, :-1]
      y = df.iloc[:, -1]
```

Check the first five observations within *X*

```
[10]:  X.head()
```

```
[10]:     R&D Spend  Administration  Marketing Spend       State
       0  165349.20        136897.80        471784.10    New York
       1  162597.70        151377.59        443898.53  California
       2  153441.51        101145.55        407934.54     Florida
       3  144372.41        118671.85        383199.62    New York
       4  142107.34         91391.77        366168.42     Florida
```

```
[11]:  X = X.values
       type(X)
```

```
[11]:  numpy.ndarray
```

Check the corresponding first five values from **Profit** column.

```
[12]:  y.head()
```

```
[12]:  0    192261.83
       1    191792.06
       2    191050.39
       3    182901.99
       4    166187.94
       Name: Profit, dtype: float64
```

```
[13]:  y = y.values
       type(y)
```

```
[13]:  numpy.ndarray
```

Encoding categorical data

```
[14]:  from sklearn.compose import ColumnTransformer
       from sklearn.preprocessing import OneHotEncoder
```

```
[15]:  ct = ColumnTransformer(transformers=[('encoder', OneHotEncoder(), [3])],␣
       ↪remainder='passthrough')
       X = np.array(ct.fit_transform(X))
```

```
[16]:  print(X[:5])
```

```
[[0.0 0.0 1.0 165349.2 136897.8 471784.1]
 [1.0 0.0 0.0 162597.7 151377.59 443898.53]
```

```
[0.0 1.0 0.0 153441.51 101145.55 407934.54]
[0.0 0.0 1.0 144372.41 118671.85 383199.62]
[0.0 1.0 0.0 142107.34 91391.77 366168.42]]
```

Splitting the dataset into training set and test set

```
[17]: from sklearn.model_selection import train_test_split
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,␣
      ↪random_state=123)
```

Training the multiple linear regression model on the training set

```
[19]: from sklearn.linear_model import LinearRegression
```

This code will create a linear regression model that fits a line to the training data, in order to make future predictions on the test data.

```
[20]: lr = LinearRegression()
      lr.fit(X_train, y_train)
```

```
[20]: LinearRegression()
```

```
[21]: theta = lr.coef_
      theta
```

```
[21]: array([-1.455e+02, -4.153e+02,  5.607e+02,  7.753e-01, -1.645e-02,
              3.627e-02])
```

```
[22]: b = lr.intercept_
      b
```

```
[22]: 48661.699896543345
```

Consider the sample `tst` as follows:

```
[23]: tst = np.array([1, 0, 0, 15e+3, 10e+2, 5e+6])
```

Predict the outcome if `tst` is the input.

```
[24]: pred = theta @ tst + b
      print('%.3f' % pred)
```

```
241495.528
```

By calling our `lr`, we get the same result:

```
[25]: lr.predict(tst.reshape(1, -1))
```

```
[25]: array([241495.528])
```

If we don't want to do the encoding of state feature by ourselves, we can invoke the previous `ct` object.

```
[26]: tst_new = [[15e+3, 10e+2, 5e+6, 'California']]
      arr = np.array(ct.transform(tst_new))
      arr
```

```
[26]: array([[1.0, 0.0, 0.0, 15000.0, 1000.0, 5000000.0]], dtype=object)
```

```
[27]: lr.predict(arr)
```

```
[27]: array([241495.528])
```

Evaluation and Visualization

Make predictions using the *X* test set and visualize the results

```
[28]: y_pred = lr.predict(X_test)
```

```
[29]: # y_test vs. _pred
      plt.scatter(y_test, y_pred, c='red')
      plt.grid()
```



```
[30]: from sklearn.metrics import mean_absolute_error, mean_squared_error,␣
        ↪mean_absolute_percentage_error
```

```
[31]: mae = mean_absolute_error(y_test, y_pred)
      mse = mean_squared_error(y_test, y_pred)
      mape = mean_absolute_percentage_error(y_test, y_pred) # relative error: *100
```

```
[32]: mae, mse, mape
```

```
[32]: (5256.526414805619, 50037959.16426007, 0.05142299211193663)
```

Multiple linear regression can be used to understand the relationship between multiple independent variables and a single dependent variable, and can be used to make predictions about the dependent variable given new data. However, it's important to note that the independent variables must be linearly related to the dependent variable in order for multiple linear regression to behave appropriately. If the relationship is non-linear, we need to use a different type of regression analysis such as polynomial regression.

Task № 2:

Using Orange Data Mining app, do the same exercise of linear regression with the predefined widgets. You can follow the rules as below:

a) import and clean the data:

- load the `50_startups.csv` file;

- generate basic statistical measures;

- create visualizations to understand data distribution;

- encode the data as demonstrated previously;

- divide the dataset into train and test sets: use 80% for training, 20% for testing.

b) apply linear regression on the training data and train it using the prepared features;

c) evaluate the model against the test test;

d) assess the model performance metrics; and

e) predict the output for the same input as in cell #23.

▼ **Remark 2**

To Launch Orange Data Mining Application, go to your already opened terminal and type `orange` *in a new tab.*

⌦ **Terminal**

student@isetbz:~$ orange

3 | *k*-NN for Classification

Student's name	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
Score /20	. . . .	. . . .	. . . .

Detailed Credits

Anticipation *(4 points)*	. . . .	. . . .	. . . .
Management *(2 points)*	. . . .	. . . .	. . . .
Testing *(7 points)*	. . . .	. . . .	. . . .
Data Logging *(3 points)*	. . . .	. . . .	. . . .
Interpretation *(4 points)*	. . . .	. . . .	. . . .

> **Motivations**
>
> ★ *k*-nearest neighbors (*k*-NN) is a simple and effective classification algorithm that is easy to understand and implement. It is based on the idea of using the class labels of the "nearest neighbors" to predict the class label of a new data point.
> ★ *k*-NN is a "lazy learner" that does not make any assumptions about the underlying data distribution, which makes it a good choice for working with complex or non-linear data. It is also robust to noise and can handle missing data. As a result, *k*-NN is often used as a baseline method for comparison with more advanced classification algorithms.

⚠ The notebook is available at `https://github.com/a-mhamdi/mlpy/` → *Codes* → *Python* → *Jupyter* → *k-nearest-neighbors.ipynb*

k-nearest neighbors (*k*-NN) is a type of instance-based learning, a method of supervised machine learning. It is used for classification and regression tasks.

In *k*-NN, the algorithm is given a labeled training dataset and a set of test data. To make a prediction for a test instance, the algorithm looks at the *k* nearest neighbors in the training dataset, based on the distance between the test instance and the training instances. The prediction is then made based on the majority class among the *k* nearest neighbors. For classification tasks, the prediction is the class with the most neighbors. For regression tasks, the prediction is the mean or median of the values of the *k* nearest neighbors.

Importing the libraries

```
[1]: import pandas as pd
```

Importing the dataset

```
[2]: df = pd.read_csv('./Datasets/Social_Network_Ads.csv')
     df.head()
```

```
[2]:    Age  EstimatedSalary  Purchased
     0   19            19000          0
     1   35            20000          0
     2   26            43000          0
     3   27            57000          0
     4   19            76000          0
```

```
[3]: X = df.iloc[:, :-1].values
     y = df.iloc[:, -1].values
```

Splitting the dataset into the Training set and Test set

```
[4]: from sklearn.model_selection import train_test_split
     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,␣
      ↪random_state=123)
```

Feature Scaling

k-NN is sensitive to the scale of the features, and it may not perform well if the features have very different scales.

```
[5]: from sklearn.preprocessing import StandardScaler
```

In order to avoid *information leakage*, it is highly important to keep in mind that only the `transform` method has to be applied on the `X_test`. (μ, σ) are of `X_train` set.

```
[6]: sc = StandardScaler()
     X_train = sc.fit_transform(X_train)
     X_test = sc.transform(X_test)
```

Training the k-NN model on the training set

```
[7]: from sklearn.neighbors import KNeighborsClassifier
```

```
[8]: clf = KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2)
```

```
[9]: clf.fit(X_train, y_train)
```

```
[9]: KNeighborsClassifier()
```

Predicting a new result

```
[10]: clf.predict(sc.transform([[30,87000]]))
```

```
[10]: array([0])
```

Predicting the test set results

```
[11]: y_pred = clf.predict(X_test)
```

Displaying the Confusion Matrix

```
[12]: from sklearn.metrics import confusion_matrix
```

```
[13]: cm = confusion_matrix(y_test, y_pred, labels=clf.classes_)
```

```
[14]: cm
```

```
[14]: array([[54,  9],
             [ 4, 33]])
```

```
[15]: from sklearn.metrics import ConfusionMatrixDisplay
```

```
[16]: ConfusionMatrixDisplay(cm, display_labels=clf.classes_).plot();
```



```
[17]: from sklearn.metrics import accuracy_score
```

```
[18]: print(f'Accuracy = {accuracy_score(y_test, y_pred):.2f}')
```

```
Accuracy = 0.87
```

https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.crosstab.html

```
[19]: pd.crosstab(y_test, y_pred, rownames=['Expected'], colnames=['Predicted'],␣
      →margins=True)
```

```
[19]: Predicted   0   1  All
      Expected
      0          54   9   63
      1           4  33   37
      All        58  42  100
```

```
[20]: from sklearn.metrics import classification_report
```

```
[21]: print(classification_report(y_test, y_pred))
```

```
                precision    recall  f1-score   support

           0        0.93      0.86      0.89        63
           1        0.79      0.89      0.84        37

    accuracy                            0.87       100
   macro avg        0.86      0.87      0.86       100
weighted avg        0.88      0.87      0.87       100
```

k-NN is a simple and effective method for classification and regression tasks, and it is easy to understand and implement. However, it can be computationally expensive to find the *k* nearest neighbors for each test instance, especially for large datasets.

> **Task № 3:**
>
> Perform the same *k*-NN experiment using the preconfigured widgets in the Orange Data Mining app. The guidelines are as follows:
>
> **a)** import and clean the data:
>
> - load the `Social_Network_Ads.csv` file;
>
> - generate basic statistical measures;
>
> - create visualizations to understand data distribution;
>
> - divide the dataset into train and test sets: use 75% for training, 25% for testing;
>
> - convert the value of the attributes to a comparable scale;
>
> **b)** apply *k*-NN on the training data and fit it using the prepared features;
>
> **c)** evaluate the model against the test test;
>
> **d)** assess the model performance metrics.

4 | K-Means for Clustering

Student's name	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
	. . . .	. . . .	. . . .
Score /20	. . . .	. . . .	. . . .

Detailed Credits

Anticipation *(4 points)*	. . . .	. . . .	. . . .
Management *(2 points)*	. . . .	. . . .	. . . .
Testing *(7 points)*	. . . .	. . . .	. . . .
Data Logging *(3 points)*	. . . .	. . . .	. . . .
Interpretation *(4 points)*	. . . .	. . . .	. . . .

> **Motivations**
>
> ★ K-means clustering is a widely used method for partitioning a dataset into a set of clusters, where each cluster consists of data points that are similar to each other. This can be useful for a variety of applications, including data compression, anomaly detection, and customer segmentation.
> ★ It can also help to identify outliers and anomalies in the data, which can be useful for identifying errors or identifying new opportunities for analysis.

The notebook is available at https://github.com/a-mhamdi/mlpy/ → *Codes* → *Python* → *Jupyter* → *k-means-clustering.ipynb*

In unsupervised learning, the algorithm is given a dataset and is asked to learn the underlying structure of the data. The goal is to find patterns or relationships in the data that can be used to group the data points into clusters or to reduce the dimensionality of the data.

Some examples of unsupervised learning algorithms include:

· *K*-Means clustering;

· Principal Component Analysis (PCA); and

· Autoencoders.

These algorithms can be used for tasks such as image compression, anomaly detection, and customer segmentation.

K-Means clustering is a method of unsupervised machine learning used to partition a dataset into *k* clusters,

where *k* is a user-specified number. The goal of *K*-Means clustering is to minimize the sum of squared distances between the points in each cluster and its centroid.

Importing the libraries

```
[1]: import pandas as pd
     import matplotlib.pyplot as plt
```

```
[2]: # Show plots in an interactive format, e.g., zooming, saving, etc
     %matplotlib inline
```

```
[3]: plt.style.use('ggplot')
```

Importing the dataset

```
[4]: df = pd.read_csv('./Datasets/Mall_Customers.csv')
```

```
[5]: df.head()
```

```
[5]:    CustomerID  Gender  Age  Annual Income (k$)  Spending Score (1-100)
     0           1    Male   19                  15                      39
     1           2    Male   21                  15                      81
     2           3  Female   20                  16                       6
     3           4  Female   23                  16                      77
     4           5  Female   31                  17                      40
```

```
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
 #   Column                  Non-Null Count  Dtype
---  ------                  --------------  -----
 0   CustomerID              200 non-null    int64
 1   Gender                  200 non-null    object
 2   Age                     200 non-null    int64
 3   Annual Income (k$)      200 non-null    int64
 4   Spending Score (1-100)  200 non-null    int64
dtypes: int64(4), object(1)
memory usage: 7.9+ KB
```

```
[7]: df.describe()
```

```
[7]:        CustomerID         Age  Annual Income (k$)  Spending Score (1-100)
     count  200.000000  200.000000          200.000000              200.000000
     mean   100.500000   38.850000           60.560000               50.200000
     std     57.879185   13.969007           26.264721               25.823522
     min      1.000000   18.000000           15.000000                1.000000
     25%     50.750000   28.750000           41.500000               34.750000
     50%    100.500000   36.000000           61.500000               50.000000
     75%    150.250000   49.000000           78.000000               73.000000
     max    200.000000   70.000000          137.000000               99.000000
```

[8]:
```
df.rename(columns={'Annual Income (k$)':'Annual Income', 'Spending Score␣
 ↪(1-100)': 'Spending Score'}, inplace=True)
```

[9]:
```
X = df.drop(columns=['CustomerID', 'Age', 'Gender']).values
X[:10, :]
```

[9]:
```
array([[15, 39],
       [15, 81],
       [16,  6],
       [16, 77],
       [17, 40],
       [17, 76],
       [18,  6],
       [18, 94],
       [19,  3],
       [19, 72]])
```

Import *K*–Means class

[10]:
```
from sklearn.cluster import KMeans
```

Training the *K*-Means model on the dataset

This code will create a *K*-Means model with 5 clusters and fit it to the data. It will then make predictions about which cluster each data point belongs to.

[12]:
```
kmeans = KMeans(n_clusters=5, init='k-means++', n_init='auto')
y_pred = kmeans.fit_predict(X)
```

[13]:
```
centers = kmeans.cluster_centers_
centers
```

[13]:
```
array([[88.2       , 17.11428571],
       [25.72727273, 79.36363636],
       [55.2962963 , 49.51851852],
       [26.30434783, 20.91304348],
       [86.53846154, 82.12820513]])
```

Visualizing the clusters

[14]:
```
fig, ax = plt.subplots()
scatter =  ax.scatter(X[:, 0], X[:, 1], c=y_pred, s=100)
legend = ax.legend(*scatter.legend_elements(), title='Clusters')
ax.add_artist(legend)
ax.scatter(centers[:, 0], centers[:, 1], c='red', s=200)
ax.set_title('Clusters of customers')
ax.set_xlabel('Annual Income')
ax.set_ylabel('Spending Score')
ax.grid()
```

Clusters of customers

Unsupervised learning can be useful when there is no labeled training data available, or when the goal is to discover patterns or relationships in the data rather than to make predictions. However, it can be more difficult to evaluate the performance of unsupervised learning algorithms, as there is no ground truth to compare the predictions to.

K-Means clustering is a fast and efficient method for clustering large datasets, and is often used as a baseline method for comparison with other clustering algorithms. However, it can be sensitive to the initial selection of centroids, and may not always find the optimal clusters if the data is not well-separated or has a non-convex shape. It is also limited to spherical clusters and may not work well for clusters with more complex shapes.

> **Task № 4:**
>
> Use the preconfigured widgets in the Orange Data Mining app to perform the same K-Means exercise. You can abide by the guidelines listed below:
>
> **a)** import and clean the data:
>
> - load the `Mall_Customers.csv` file;
>
> - generate basic statistical measures;
>
> - create visualizations to understand data distribution;
>
> - encode the data as demonstrated previously;
>
> - divide the dataset into train and test sets: use 80% for training, 20% for testing.
>
> **b)** apply K-Means on the training data and fit it using the prepared features;
>
> **c)** evaluate the model against the test test;
>
> **d)** assess the model performance metrics.

A. Mhamdi

5 | Binary Classifier using ANN

Student's name	. . . .	. . . .	. . . .
Score /20	. . . .	. . . .	. . . .

Detailed Credits

Anticipation (4 points)	. . . .	. . . .	. . . .
Management (2 points)	. . . .	. . . .	. . . .
Testing (7 points)	. . . .	. . . .	. . . .
Data Logging (3 points)	. . . .	. . . .	. . . .
Interpretation (4 points)	. . . .	. . . .	. . . .

Motivations

★ For binary classification tasks, such as predicting a binary outcome (such as "yes" or "no") based on input data, artificial neural networks (ANNs) are an effective tool. ANNs are well-suited for jobs with a complex underlying structure or a large number of features because they can learn intricate correlations between the input data and the output labels.

★ ANNs are highly flexible and can be trained on a wide range of data types, including continuous and categorical variables. They can also handle missing values and handle large amounts of data efficiently. This makes them a good choice for tasks where the data is noisy or high-dimensional.

The notebook is available at `https://github.com/a-mhamdi/mlpy/` → *Codes* → *Python* → *Jupyter* → *artificial-neural-network.ipynb*

Artificial neural networks (ANN) are commonly used for classification tasks because they are able to learn complex relationships between the input features and the target class. They are particularly useful when the relationship is non-linear, as they are able to learn and model the inputs-outputs mapping using multiple hidden layers of interconnected neurons.

ANN are also able to handle large amounts of data and can learn from it without being explicitly programmed with a set of rules or a decision tree. This allows them to be very flexible and adaptable, and makes them well-suited for tasks that are difficult to define using traditional programming techniques.

There are several advantages to using neural networks for classification tasks:

1. They are able to learn complex relationships between the input features and the target class;
2. They are able to handle large amounts of data;

3. They can learn from unstructured data;
4. They are flexible and adaptable;
5. They can be trained to perform well on a wide range of classification tasks.

Importing the libraries

```python
[1]: import numpy as np
     import matplotlib.pyplot as plt
     import pandas as pd
```

```python
[2]: np.set_printoptions(precision=2)
```

Importing the dataset

```python
[3]: df = pd.read_csv("./Datasets/Churn_Modelling.csv")
```

```python
[4]: df = df.dropna(how="any", axis=0)
```

```python
[5]: df.head()
```

```
[5]:    RowNumber  CustomerId   Surname  CreditScore Geography  Gender  Age  \
    0          1    15634602  Hargrave          619    France  Female   42
    1          2    15647311      Hill          608     Spain  Female   41
    2          3    15619304      Onio          502    France  Female   42
    3          4    15701354      Boni          699    France  Female   39
    4          5    15737888  Mitchell          850     Spain  Female   43

       Tenure     Balance  NumOfProducts  HasCrCard  IsActiveMember  \
    0       2        0.00              1          1               1
    1       1    83807.86              1          0               1
    2       8   159660.80              3          1               0
    3       1        0.00              2          0               0
    4       2   125510.82              1          1               1

       EstimatedSalary  Exited
    0         101348.88       1
    1         112542.58       0
    2         113931.57       1
    3          93826.63       0
    4          79084.10       0
```

```python
[6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype
---  ------            --------------  -----
 0   RowNumber         10000 non-null  int64
 1   CustomerId        10000 non-null  int64
 2   Surname           10000 non-null  object
 3   CreditScore       10000 non-null  int64
```

```
4    Geography         10000 non-null   object
5    Gender            10000 non-null   object
6    Age               10000 non-null   int64
7    Tenure            10000 non-null   int64
8    Balance           10000 non-null   float64
9    NumOfProducts     10000 non-null   int64
10   HasCrCard         10000 non-null   int64
11   IsActiveMember    10000 non-null   int64
12   EstimatedSalary   10000 non-null   float64
13   Exited            10000 non-null   int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

[7]: `df.describe()`

[7]:

	RowNumber	CustomerId	CreditScore	Age	Tenure
count	10000.00000	1.000000e+04	10000.000000	10000.000000	10000.000000
mean	5000.50000	1.569094e+07	650.528800	38.921800	5.012800
std	2886.89568	7.193619e+04	96.653299	10.487806	2.892174
min	1.00000	1.556570e+07	350.000000	18.000000	0.000000
25%	2500.75000	1.562853e+07	584.000000	32.000000	3.000000
50%	5000.50000	1.569074e+07	652.000000	37.000000	5.000000
75%	7500.25000	1.575323e+07	718.000000	44.000000	7.000000
max	10000.00000	1.581569e+07	850.000000	92.000000	10.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember
count	10000.000000	10000.000000	10000.00000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

[8]:
```
X = df.iloc[:, 3:-1].values
y = df.iloc[:, -1].values
```

Data preprocessing

```
[9]:  from sklearn.preprocessing import LabelEncoder, OneHotEncoder
```

```
[10]: le = LabelEncoder()
      ohe = OneHotEncoder()
```

```
[11]: X[:, 2] = le.fit_transform(X[:, 2])
```

```
[12]: from sklearn.compose import ColumnTransformer
```

```
[13]: ct = ColumnTransformer([("ohe", ohe, [1])], remainder='passthrough')
      X = np.array(ct.fit_transform(X))
```

```
[14]: X[:5, :]
```

```
[14]: array([[1.0, 0.0, 0.0, 619, 0, 42, 2, 0.0, 1, 1, 1, 101348.88],
             [0.0, 0.0, 1.0, 608, 0, 41, 1, 83807.86, 1, 0, 1, 112542.58],
             [1.0, 0.0, 0.0, 502, 0, 42, 8, 159660.8, 3, 1, 0, 113931.57],
             [1.0, 0.0, 0.0, 699, 0, 39, 1, 0.0, 2, 0, 0, 93826.63],
             [0.0, 0.0, 1.0, 850, 0, 43, 2, 125510.82, 1, 1, 1, 79084.1]],
            dtype=object)
```

```
[15]: X = np.asarray(X, dtype=np.float64)
```

```
[16]: X[:5, :]
```

```
[16]: array([[1.00e+00, 0.00e+00, 0.00e+00, 6.19e+02, 0.00e+00, 4.20e+01,
              2.00e+00, 0.00e+00, 1.00e+00, 1.00e+00, 1.00e+00, 1.01e+05],
             [0.00e+00, 0.00e+00, 1.00e+00, 6.08e+02, 0.00e+00, 4.10e+01,
              1.00e+00, 8.38e+04, 1.00e+00, 0.00e+00, 1.00e+00, 1.13e+05],
             [1.00e+00, 0.00e+00, 0.00e+00, 5.02e+02, 0.00e+00, 4.20e+01,
              8.00e+00, 1.60e+05, 3.00e+00, 1.00e+00, 0.00e+00, 1.14e+05],
             [1.00e+00, 0.00e+00, 0.00e+00, 6.99e+02, 0.00e+00, 3.90e+01,
              1.00e+00, 0.00e+00, 2.00e+00, 0.00e+00, 0.00e+00, 9.38e+04],
             [0.00e+00, 0.00e+00, 1.00e+00, 8.50e+02, 0.00e+00, 4.30e+01,
              2.00e+00, 1.26e+05, 1.00e+00, 1.00e+00, 1.00e+00, 7.91e+04]])
```

```
[17]: from sklearn.model_selection import train_test_split
```

```
[18]: X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=.8,␣
      ↪random_state=123)
```

```
[19]: from sklearn.preprocessing import MinMaxScaler
```

```
[20]: sc = MinMaxScaler()
```

```
[21]: X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

```
[22]: print(X_train[:5, :])
```

```
[[1.    0.    0.    0.75 1.    0.16 1.    0.    0.33 1.    1.    0.27]
 [1.    0.    0.    0.51 0.    0.28 1.    0.    0.67 1.    0.    0.66]
 [0.    0.    1.    0.87 1.    0.23 0.3  0.    0.33 0.    0.    0.41]
 [1.    0.    0.    0.69 1.    0.3  0.9  0.    0.33 1.    0.    0.2 ]
 [1.    0.    0.    0.71 1.    0.2  0.3  0.58 0.33 1.    0.    0.57]]
```

Build the classifier `clf`

```python
[23]: from keras.models import Sequential
      from keras.layers import Dense
```

```python
[24]: clf = Sequential()
      ndim = X_train.shape[1]
      clf.add(Dense(units=8, activation="relu", input_dim=ndim))
      clf.add(Dense(units=4, activation="relu"))
      clf.add(Dense(units=4, activation="relu"))
      clf.add(Dense(units=1, activation="sigmoid"))
```

Insights about `clf`

```python
[25]: clf.summary()
```

```
Model: "sequential"

_____
 Layer (type)                Output Shape              Param #
=================================================================
 dense (Dense)               (None, 8)                 104

 dense_1 (Dense)             (None, 4)                 36

 dense_2 (Dense)             (None, 4)                 20

 dense_3 (Dense)             (None, 1)                 5

=================================================================
Total params: 165
Trainable params: 165
Non-trainable params: 0
_____
```

Compile `clf`

```python
[26]: clf.compile(optimizer="adam",
                  loss="binary_crossentropy",
                  metrics=['Accuracy', 'Precision', 'Recall'])
```

Train and evaluate `clf`

```python
[27]: clf.fit(X_train, y_train, batch_size=16, epochs=32);
```

```
Epoch 1/32
500/500 [==============================] - 1s 1ms/step - loss: 0.5106 -
```

```
Accuracy: 0.7904 - precision: 0.3533 - recall: 0.0400

...

Epoch 32/32
500/500 [==============================] - 1s 1ms/step - loss: 0.3545 -
Accuracy: 0.8558 - precision: 0.7492 - recall: 0.4344
```

▼ Remark 3

The previous cell's output has been shortened. Check the entire `.ipynb` *notebook.*

[28]:
```python
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
```

[29]:
```python
y_pred = clf.predict(X_test)
y_pred = ( y_pred > .5 ).astype(int)
```

```
63/63 [==============================] - 0s 773us/step
```

Print the confusion matrix

[30]:
```python
cm = confusion_matrix(y_test, y_pred)
```

[31]:
```python
ConfusionMatrixDisplay(cm).plot();
```


[32]:
```python
print(y_test.shape)
```

```
(2000,)
```

[33]:
```python
print(y_pred.shape)
y_pred = y_pred.reshape(len(y_pred),)
print(y_pred.shape)
```

```
(2000, 1)
(2000,)
```

```
[34]: pd.crosstab(y_test, y_pred, rownames=["Expected"], colnames=["Predicted"],␣
      ↪margins=True)
```

```
[34]: Predicted      0    1    All
      Expected
      0           1529   57  1586
      1            217  197   414
      All         1746  254  2000
```

```
[35]: y_test = y_test.reshape(len(y_test), 1)
      y_pred = y_pred.reshape(len(y_pred), 1)
      print(np.concatenate((y_test[:10], y_pred[:10]), axis=1))
```

```
[[0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [0 0]
 [1 0]
 [1 1]
 [0 0]
 [0 0]]
```

```
[36]: from sklearn.metrics import classification_report
```

```
[37]: print(classification_report(y_test, y_pred))
```

```
               precision    recall  f1-score   support

           0        0.88      0.96      0.92      1586
           1        0.78      0.48      0.59       414

    accuracy                            0.86      2000
   macro avg        0.83      0.72      0.75      2000
weighted avg        0.85      0.86      0.85      2000
```

It is important to note that neural networks can be more computationally intensive to train and may require more data and more time to achieve good performance, compared to some other classification algorithms. Additionally, they can be more difficult to interpret and understand, as they learn patterns in the data through the weights and biases of the network rather than through explicit rules.

Task № 5:

Use the preloaded widgets in the Orange Data Mining app to complete the identical artificial network exercise. The guidelines are:

a) import and clean the data:

- load the `Churn_Modelling.csv` file;
- generate basic statistical measures;
- create visualizations to understand data distribution;

- divide the dataset into train and test sets: use 80% for training, 20% for testing.

- scale the features value to a similar scale;

b) apply ANN on the training data and fit it using the prepared features;

c) evaluate the model against the test test;

d) assess the model performance metrics.

The overall scope of this manual is to introduce **Machine Learning**, through some numeric simulations, to the students at the department of **Electrical Engineering**.

The topics discussed in this manuscript are as follow:

① Getting started with *Python*

② Linear Regression

③ Classification

④ Clustering

⑤ ANN

Python; *Orange Data Mining, Jupyter*; *Marimo*; *NumPy*; *Matplotlib*; *Scikit-learn*; *Keras*; machine learning; linear regression; classification; clustering; ann.